

- **Types of ML:** unsupervised(learn a model of data), supervised(learn mapping from input to target), reinforcement(learn from mistakes), imitation(type of reinforcement)
- **Linear Regression:** supervised,  $X^{N,1+D}W^{1+D,1} = y^{N,1}$ 
  - Minimize SSE(sum squared error)
  - With SSE,  $W = (X^T X)^{-1} X^T t$
- **Gradient Descent**
  - $w(t+1) = w(t) - \alpha(\partial J / \partial w)$
  - Batch descent:  $w_i = w_i + \alpha \frac{1}{N} \sum_{n=1}^N \delta^n x_i^n$
  - SGD: update weight 1by1, needs shuffle
- **Perceptron**
  - Outputs 1 or 0
  - $w_i = w_i + \alpha(t^n - y^n)x_i^n$  where  $\delta^n = (t^n - y^n)$
  - Can learn whatever it can compute, but slow
- **Maximum Likelihood Objective Function**
  - Adjust your parameters to maximize likelihood of data distribution you see
  - SSE assumes Gaussian distribution and cross entropy assumes Bernoulli/multinomial distribution
  - We can cluster datapoints by pushing same category together and pushing different one apart—supervised clustering(we know data's category)
    - Siamese networks: two identical NN
    - Loss function:  $L(W, Y, \vec{x}_1, \vec{x}_2) = (1 - Y)\frac{1}{2}(Dw)^2 + (Y)\frac{1}{2}\{\max(0, m - Dw)\}^2$  where D is distance between networks' output, m is margin, Y is indicator for different pair
- **Backprop**
  - $w_{ij} = w_{ij} + \alpha \delta_j z_i$  where  $\delta_j = t_j - y_j$  for output unit and  $g'(a_j) \sum_k \delta_k w_{jk}$  for hidden unit
- **Generalization**
  - More data/augment existing data
  - Regularization: minimize  $J = E + \lambda C$  to penalize too complex model
    - $C = |W|(L1); C = ||W||_2^2(L2)$
  - Dropout: randomly turn off hidden units when training
  - Early Stopping: leave a holdout set from training data, watch the error on it and stop training if it starts to rise
  - Add Gaussian noise into inputs/model/outputs
- **Tricks for Training**
  - SGD and minibatch: faster to converge, better generalization, adaptive, take advantage of dataset redundancy, need shuffling
  - Normalize data
    - All positive inputs make weight changes all + or -
    - Correlated inputs: redundant information
    - Different scale: largest weight change for large inputs, bad if two inputs are equally important
  - PCA: shifts mean of input var to 0, decorrelates inputs, dimensionality reduction
  - Z-scoring: shift mean to 0, make variables same size, no decorrelate, no dimension reduction
  - Activation function
    - Sigmoid: all positive output, bad. Use  $f(x) = 1.7159 \tanh(0.667x)$  instead
    - ReLU:  $\max(0, x)$  only non-negative output, input issue not as big of concern, also can use batch normalization
  - Weight initialization
    - Init 0, if sigmoid, delta of hidden will be the same, hidden units all compute same feature. If tanh, all weights stay zero.
    - Want weighted sum of inputs to be 0 mean and unit deviation (linear range of sigmoid) because gradients will be largest and network learn any linear part before non-linear part
    - $sd(a_j) = \sqrt{\text{var}(XW)} = \sqrt{\text{var}(W)} = \sqrt{\text{sum}_i(w_{ij}^2)}$

- Fan-in: number of inputs to unit j, fan-out: number of outputs
- Let fan-in be m, init weights to mean 0, sd  $1/(m)^{1/2}$
- ReLU

Xavier initialization (Glorot & Bengio, 2010):

- Where  $n_j$  and  $n_{j+1}$  are the fan-in and fan-out of a unit, respectively)  $W \sim U\left[-\frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}}, \frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}}\right]$  (16)

Better: Kaiming initialization (He, et al., 2015):

$$W \sim \mathcal{N}\left(0, \sqrt{\frac{2}{n_i}}\right), \text{ where } n_i \text{ is the fan-in, and}$$

Here,  $\mathcal{N}$  denotes a normal distribution

- Batch normalization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1...m}\}$ ;  
Parameters to be learned:  $\gamma, \beta$   
**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

- Treated as a layer between output of previous layer the input of next layer
- Operation is differentiable so can backprop
- Convergence with a multi-layer, non-linear network indicates that the network has found a local minima along the error surface
- Momentum
  - Motivation: – Move quickly in directions with small but consistent gradients. – Move slowly in directions with big but inconsistent gradients.
  - Nesterov: update w with accumulated gradient ,calculate new gradient then make correction
- Adaptive learning rate
  - rprop for full batch and rmsprop for minibatch
- **Convnets**
  - Why convnet: images are huge!
  - Convnets use four principles:
    - Pixels depend on nearby pixels (locality): Hence, small receptive field
    - The statistics of visual inputs are invariant across images (stationary statistics): Hence, replicate receptive fields across images
    - Objects don't change identity based on location: (translation invariance): Spatial pooling
    - Objects are made of parts(compositionality): receptive fields get larger deeper in the net
  - Components: convolution(kernel dot inputs)--nonlinearity(rectified linear)--pooling
    - Convolute filter (3x3) computed on whole image(*stationary statistics*)
    - Nonlinearity: applied per pixel (ReLU)

- Spatial Pooling: max is best, avg/sum less. Pooling provides larger receptive field and invariance to small transformations
- Batch normalization: normalize layer inputs to eliminate covariance shift caused by changing previous layer parameters
- Deep network:
  - Deeper is better: more abstract feature, larger reception field
  - To pass gradients effectively: deep supervision, skip connections, etc.
  - Reuse pretrained networks and modify for new tasks
- Translation invariance(good, built into the network); scale invariance(good, learned from data, can be built in); rotational invariance(bad, can be built in)
- Dilation: increase receptive field without losing resolution. Pooling and stride also increase receptive field but reduce resolution
- Network in Network: 1x1 kernel to add nonlinearity
- Global Average Pooling: use feature map's average as output instead of softmax. Can visualize activation to see what network is looking at
- GoogleNet: multiple kernel sizes at same layer: look at different scales; 1x1 kernel for reducing dimension and number of computation
- ResNet: introduce pass through between layers to help gradient propagation
- Visualize Features: deconvolution(zero out all features except the max and propagate back); gradient descent back(optimize input to maximize particular output by taking gradient of output and backprop)
- Adversarial: calculate gradient at the output and climb up
- Training tips: use SGD/mini-batch; Adam optimizer is popular; start with big learning rate and anneal it; use Nesterov momentum; examine the data and normalization; measure both training and validation error; test with small dataset before running in full.
- Signs of good training: hidden units are sparse across samples and across features; learned filters exhibit structure and are uncorrelated.
- Problems & Fix:
  - Training diverges: learning rate too high
  - Bad accuracy: make the network bigger/deeper; visualize feature and fix optimization
  - Training is slow: use matrix operation not for loop

## ● Map time into:

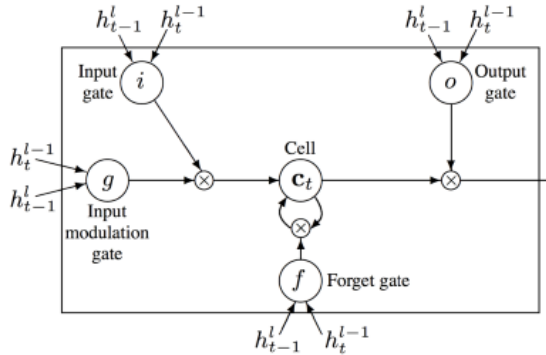
- Space: NETTalk, Transformer
  - Autoregressive models Predict the next term in a sequence from a fixed number of previous terms using "delay taps".
  - Feed-forward models generalize autoregressive models by using 1+ layers of hidden units
  - Work well for small problems
  - Difficulty dealing with variable input sizes
  - But transformer changed all this
- State: RNN, LSTM
  - Use activation memory to process input based on previous ones

## ● RNN

- Recurrent network can be unrolled in time and becomes feed-forward
- BPTT: compute gradient, propagate backward in time, and compute average change between time steps for a weight
- Problem of BPTT: backprop is linear, so gradient can vanish/explode. Hard to train long network(solved by LSTM)

## ● LSTM

- Three gates: write(input), keep(info stay in the cell through time), read(output)



$$\begin{pmatrix} i \\ f \\ o \end{pmatrix} = \begin{pmatrix} \text{sigm} \\ \text{sigm} \\ \text{tanh} \end{pmatrix} T_{2n,4n} \left( D(h_{t-1}^{l-1}) \right)$$

$$c_t^l = f \odot c_{t-1}^l + i \odot g$$

$$h_t^l = o \odot \tanh(c_t^l)$$

## ● RNN generation


- N to 1: sentiment classification; 1 to N: image captioning; N to N(after): machine translation; N to N(sync): label frames of video
- Feed previous output into input
- Combined with conv to generate image caption, can include original feature to overcome longer timesteps

## ● Transformers

- Feedforward, same weights for all inputs, parallelizable: fast to train
- A tower of encoder/decoder contains feedforward layers and attention
- Encoder has self-attention: it can look at multiple inputs
- Attention: tower communicate with other towers through attention network
  - KQV: key(what I have), query(what I'm looking for), value(what I give when other towers ask)
  - Output: Softmax(KQ)\*V (softmax of all towers' key times tower's own query forms interest vector, then times all towers' value)
- Image transformers: same architecture(encoders), no convolution, taking patches of image with positional encoding Residual: preserve gradient

## ● Neural Turing Machine

- Neural net can learn to program but slow and difficult to adapt
- NTM: add a structured memory to neural controller to read and write
- Neural controller can be feedforward or recurrent(better)



- K: key; beta: gain on content match; g: switch between content/location addressing; s: shift the address; gamma: gain on softmax address

$$w_t^c(i) \leftarrow \frac{\exp(\beta_t K[k_t, M_t(i)])}{\sum_j \exp(\beta_t K[k_t, M_t(j)])}$$

$$w_t^g \leftarrow g_t w_t^c + (1 - g_t) w_{t-1}$$

$$\tilde{w}_t(i) \leftarrow \sum_{j=0}^{N-1} w_t^g(j) s_t(i - j)$$

$$w_t(i) \leftarrow \frac{\tilde{w}_t(i)^{\gamma_t}}{\sum_j \tilde{w}_t(j)^{\gamma_t}}$$

- Can learn to program: Copy, repeated copy, associative recall, priority sort

## ● Reinforcement Learning

- An agent act in an environment, which changes states and gives the agent an reward. Goal is to maximize reward
- Agent learns a policy: given state, output a probability function of actions
- Policy gradient: sample the softmax policy distribution at each step; use the sample as teacher; compute the weight change and keep running average; multiply the weight changes with reward sign when game ends; update the network
- Discount rate: exponentially reduce potential future rewards(0-1, the bigger the more far-sighted)
- Markov property: all the agent needs to know at any time is the current state(which usually isn't true)
- Model based learning: the agent either has available to it, or it learns, a model of its environment. Most RLs are model-free
- Model: probability distribution of possible outcome states given current state and actions
- Value of state: expected reward in this state given policy
- Q learning: update current state's Q value based on Q values of states you get to. Overtime the Q values approximate final reward
- Minimax: a player maximizing own reward while the opponent minimizing its reward. Tree's branching factor equals to number of legal moves, cannot use for larger games
- Temporal Difference Gammon

$$w_{t+1} - w_t = \alpha(Y_{t+1} - Y_t) \sum_{k=1}^t \lambda^{t-k} \nabla_{w_t} Y_k$$

- The weights were updated using a temporal difference rule on every move. That is, the current estimate of the board value was updated to be closer to the next board's value
- lambda: how far error feeds back(stop when 0, arbitrarily far when 1)

## ● Alpha-go:


- Two policy networks with supervised training. Shallow for rollout, deep for training: play against younger selves with reinforcement learning
- A value network predicts winner given states: supervised training
- Monte Carlo Tree Search:
  - Traverse the tree with depth L
  - Selection actions based on Q+upper confidence bound
  - The upper confidence bound increases as a move is not tried
  - Expand the tree based on selection
  - Evaluate nodes by average of shallow value network and value network, propagate value up the tree
  - Back up the tree for future use

## ● Alpha-zero:


- One network for both value and policy
- Simplified MCTS(selection, expand & evaluate, backup)
- Use resnet with BN

## ● GAN

- Unsupervised learning to generate data from a distribution
- Recall autoencoder, which is trained to replicate the input and learns representations in the process



- Random noise
- GAN has only decoder(generator) and an adversary(discriminator)
- Discriminator learns to identify fake image from real, and generator tries to fool the discriminator
- Discriminator trains by supervised label, generator trains by gradient ascent on discriminator
- Conditional GAN: feed generator with images instead of noise
- Cycle GAN: translate domain



- ReCycle GAN: map videos to videos
  - Use Unet-based predictor to predict next frame in target domain, then translate back into original domain

## ● SimCLR: self supervision

- In a batch, augment a data to get positive pairs and use other images in the batch as negative pairs
- Add an extra layer on final output and minimize its outputs on positive pairs. Then throw away the final layer after training is complete
  - Rationale: by mapping augmented data exactly the same, some information is lost
- Use normalize temperature-scales cross entropy loss
  - Minimize distance of positive pair and vice versa
- Randomized data augmentation